Carnegie Mellon
CyLab
CONFIDENCE FOR A NETWORKED WORLD

Carnegie Mellon CyLab
4616 HENRY STREET
PITTSBURGH, PA 15213
PH: (412) 268-7195
FX: (412) 268-7196
www.cylab.cmu.edu

# SCION:

## Scalability, Control and Isolation On Next-Generation Networks

**Xin Zhang, Hsu-Chun Hsiao, Geoff Hasker, Haowen Chan, Adrian Perrig, David Andersen**

# Reasons for Clean-Slate Design

- Someone may just want to deploy a new Internet ☺
  - ✓ Possible for specialized high-reliability networks, e.g., smart grid
  - ✓ We need to have a design ready
- Even if we want to evolve current Internet, we need to have a goal, know how good a network could be

**The question is not: why deploy a new Internet?**
**But: why are we still putting up with the current Internet?**

# After years of patching, the Internet is still neither Reliable nor Secure!

Feb 2008: Pakistani ISP hijacks YouTube prefix

Apr 2010: A Chinese ISP inserts fake routes affecting thousands of US networks.

Nov 2010: 10% of Internet traffic 'hijacked' to Chinese servers due to DNS Tampering.

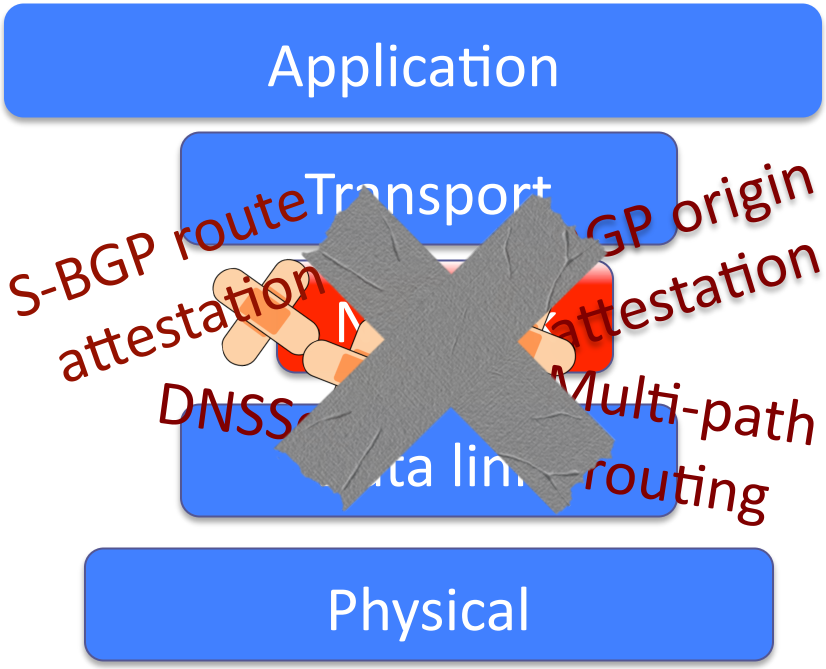

- ❖ Fixes to date – ad hoc, patches
- ❖ Inconvenient truths
  - ✧ S-BGP: delayed convergence
  - ✧ Global PKI: single root of trust

# Limitations of the Current Internet

- Destination or ISP have no control over inbound paths

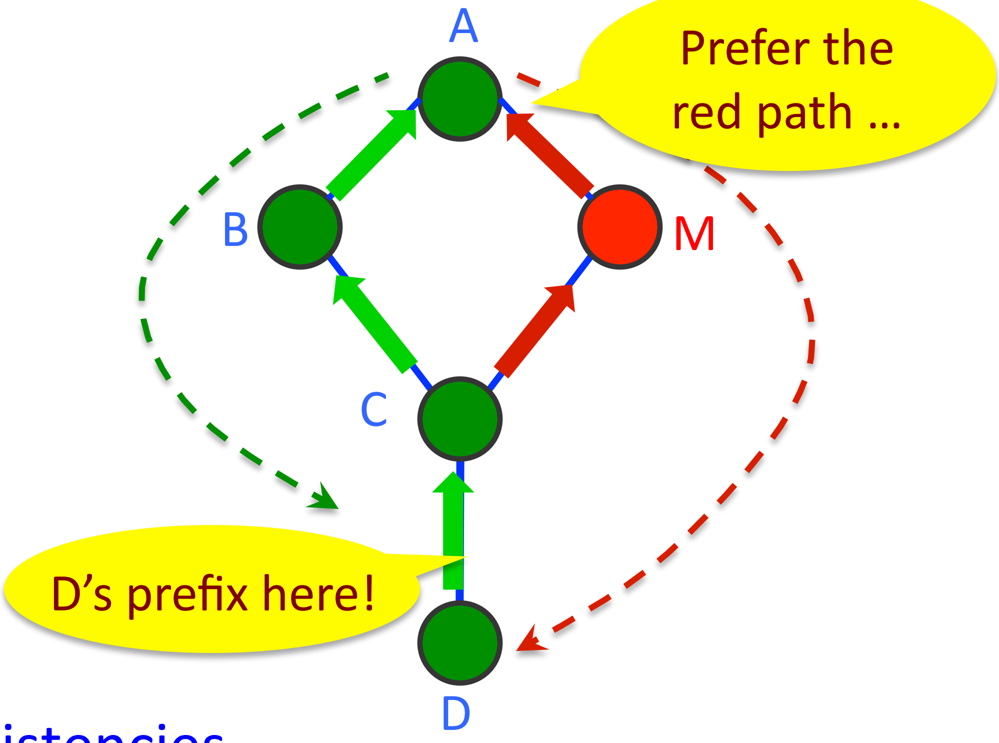

- Route inconsistencies
  - Forwarding state may be different from announced state

# Limitations of the Current Internet (cont’d)

- Lack of routing isolation
  - A failure/attack can have global effects
  - Global visibility of paths is not scalable
- Slow convergence / route oscillation
- Large routing tables
  - Multi-homing / flat namespaces prevent aggregation
- Lack of route freshness
  - Current (S-)BGP cannot prevent replay of old paths

Note that these issues are fundamental to (S)-BGP!

# Wish List (1): Isolation

- ❖ Isolation of attacks
- ❖ Scalable and reliable routing updates
- ❖ Operate with mutually distrusting entities without a global single root of trust: enforceable accountability

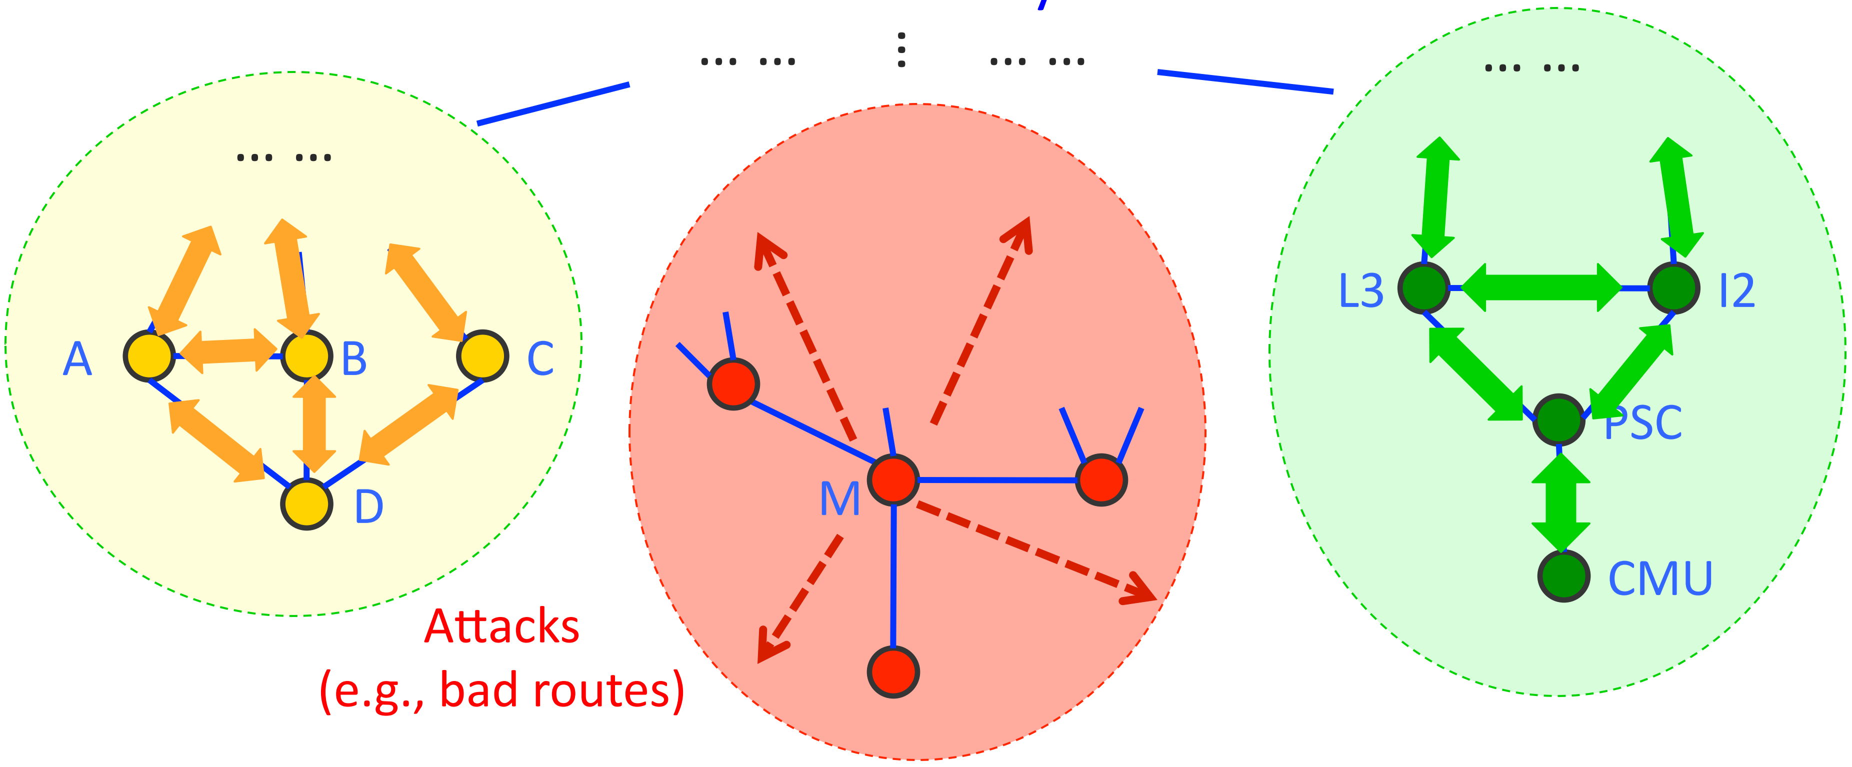

Attacks (e.g., bad routes)

# Wish List (2): Balanced Control

- Transit ISPs, source and destination all need path control

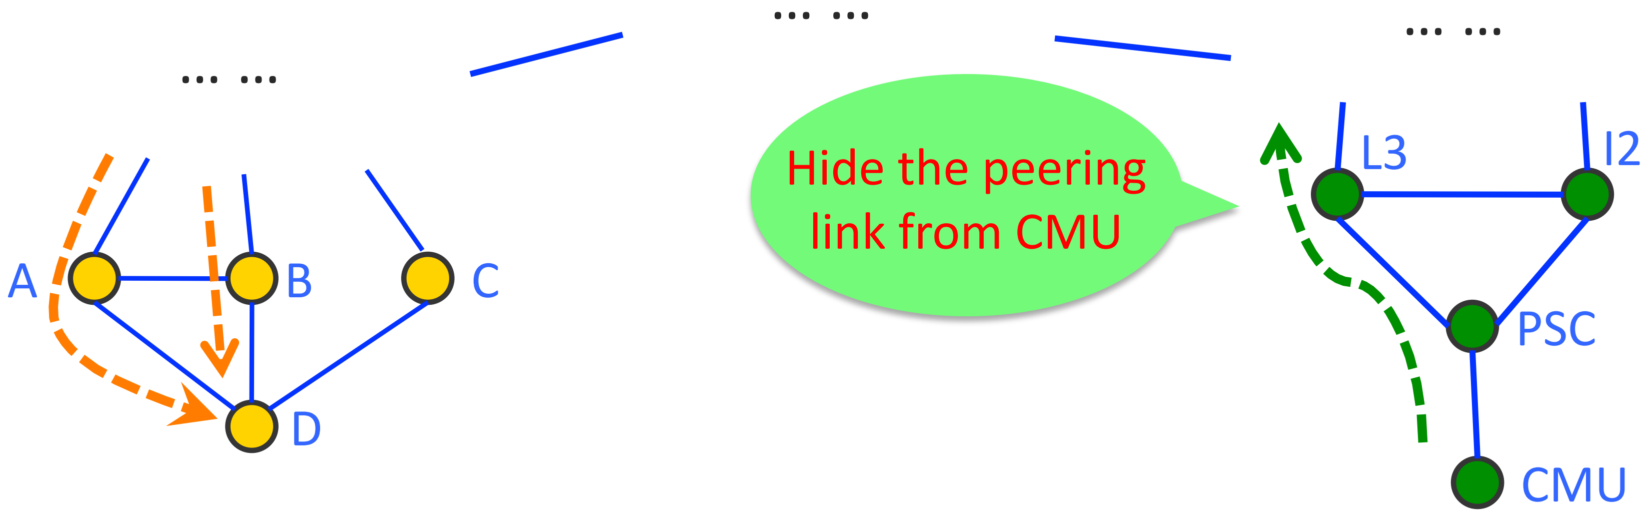

# Wish List (3): Explicit Trust

- Know who needs to be trusted
  - Absence of consistency in BGP prevents knowing exactly who needs to be trusted

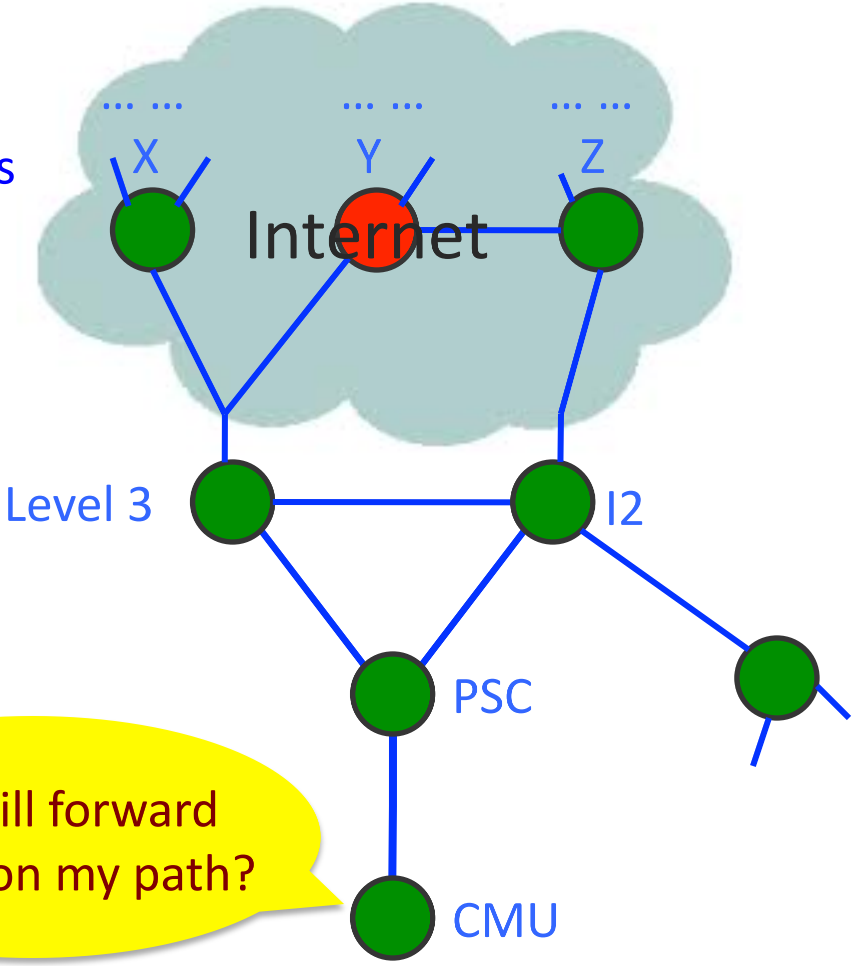

Who will forward packets on my path?

# SCION Architectural Goals

- High availability, even for networks with malicious parties
- Explicit trust for network operations
- Minimal TCB: limit number of entities that need to be trusted for any operation
  - Strong isolation from untrusted parties
- Operate with mutually distrusting entities
  - No single root of trust
- Enable route control for ISPs, receivers, senders
- Simplicity, efficiency, flexibility, and scalability

# SCION Architecture Overview

- Trust domain (TD)s
  - Isolation and scalability
- Path construction
  - Path construction beacons (PCBs)
- Path resolution
  - Control
  - Explicit trust
- Route joining (shortcuts)
  - Efficiency, flexibility

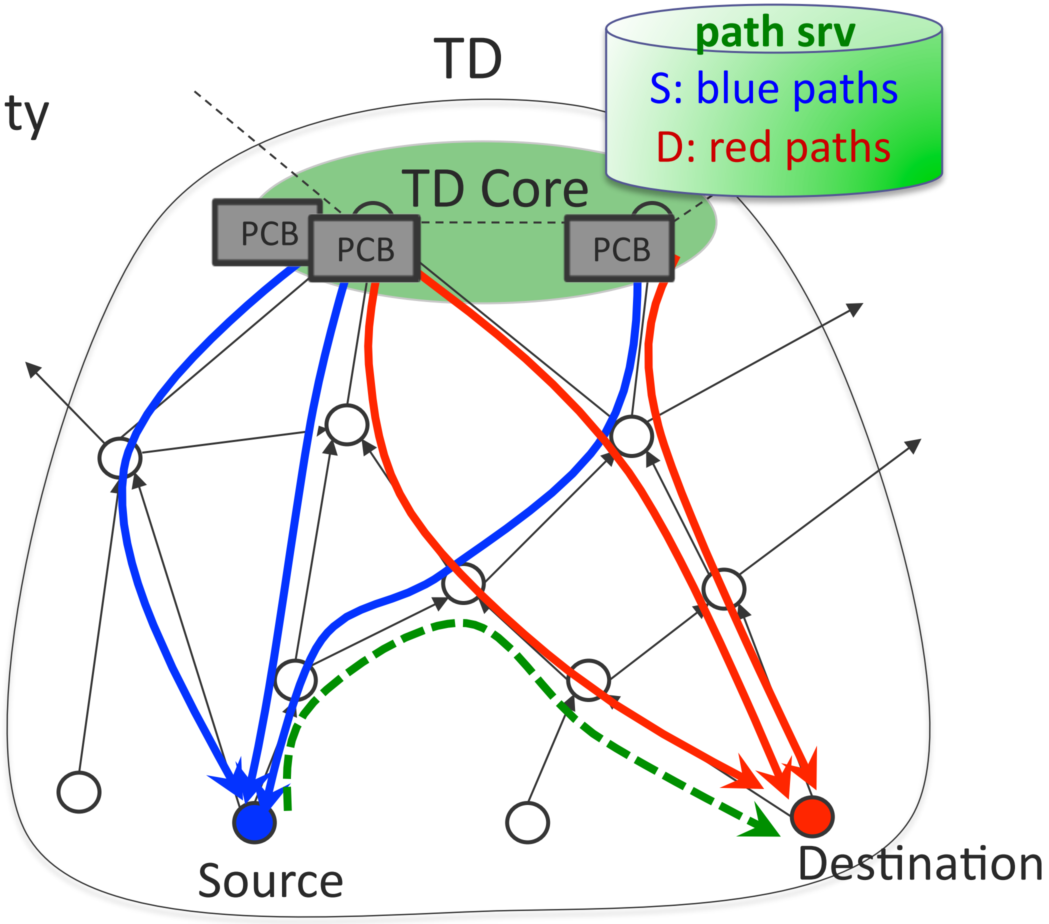

# Trust Domain Decomposition

- Global set of TD (Trust Domains)
  - ✓ Map to geographic, political, legal boundaries
- TD Core: set of top-tier ISPs that manage TD
  - ✓ Route to other TDs
  - ✓ Initiate path construction beacons
  - ✓ Manage Address and Path Translation Servers
  - ✓ Handle TD membership
  - ✓ Root of trust for TD: manage root key and certificates
- AD is atomic failure unit, contiguous/autonomous domain
  - ✓ Transit AD or endpoint AD

# Path Construction

Goal: each endpoint learns multiple verifiable paths to its core

- Discovering paths via Path Construction Beacons (PCBs)
  - ✓ TD Core periodically initiates PCBs
  - ✓ Providers advertise upstream topology to peering and customer ADs
- ADs perform the following operations
  - ✓ Collect PCBs
  - ✓ For each neighbor AD, select which *k* PCBs to forward
  - ✓ Update cryptographic information in PCBs
- Endpoint AD will receive up to *k* PCBs from each upstream AD, and select *k* down-paths and up-paths

# Path Construction

△: interface ●: Opaque field ☾: expiration time ♥: signature

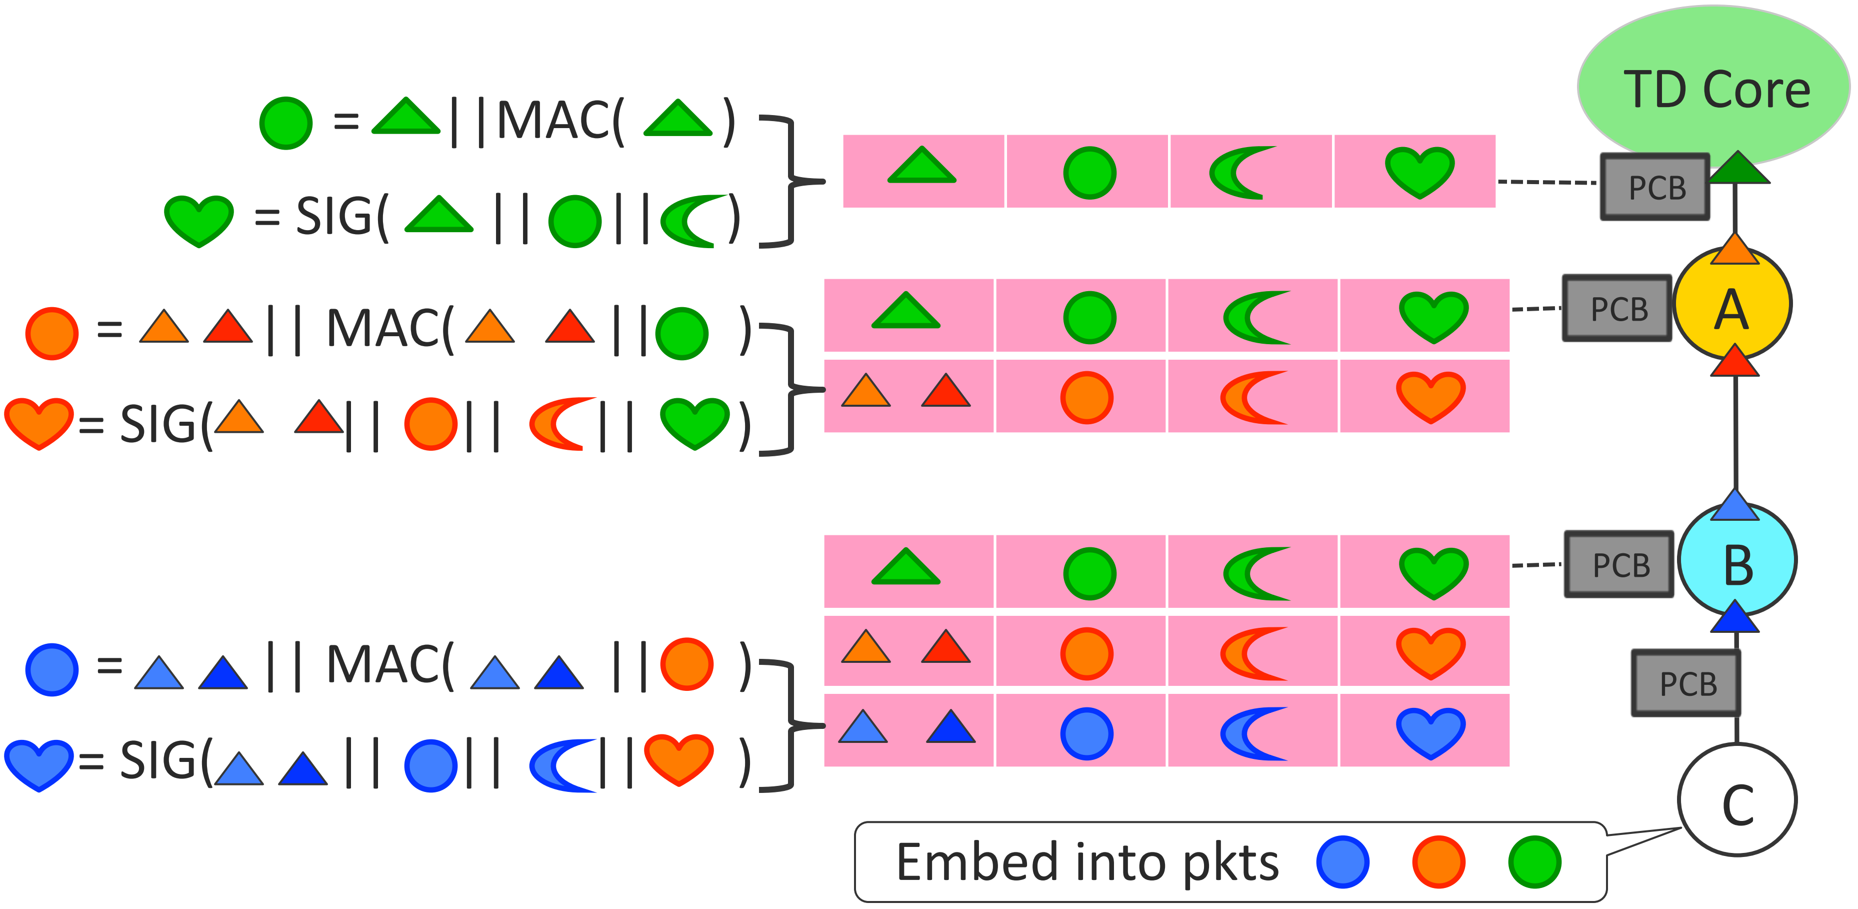

# Path Construction

*Interfaces: I(i) = previous-hop interfaces || local interfaces*

*Opaque field: O(i) = local interfaces || MAC over local interfaces and O(i-1)*

*Signature: Σ(i) = sign over I(i), T(i), O(i), and Σ(i-1), with cert of pub key*

*C*? – One PCB per neighbor

C→E: $I(C)$: $I(A)$|| {$C1$, $C4$}

$O(C)$ : {$C1$, $C4$} || $\mathrm{MAC}_{\mathrm{Ka}}$( {$C1$, $C4$} || $O(A)$ )

Σ(C): Sign( $I(C)$ || $T(C)$ || $O(C)$ ||Σ(A) )

Also include peering link!

$I_{C,D}$(C): {$C4$,$C2$} || TD || $\mathrm{AID}_{\mathrm{D}}$

$O_{C,D}$(C): {$C4$, $C2$} ||$\mathrm{MAC}_{\mathrm{Kc}}$( {$C4$, $C2$} )

$\Sigma_{C,D}$(C): Sign( $I_{C,D}$(C) || $T_{C,D}$(C) || $O_{C,D}$(C) || $O(C)$ )

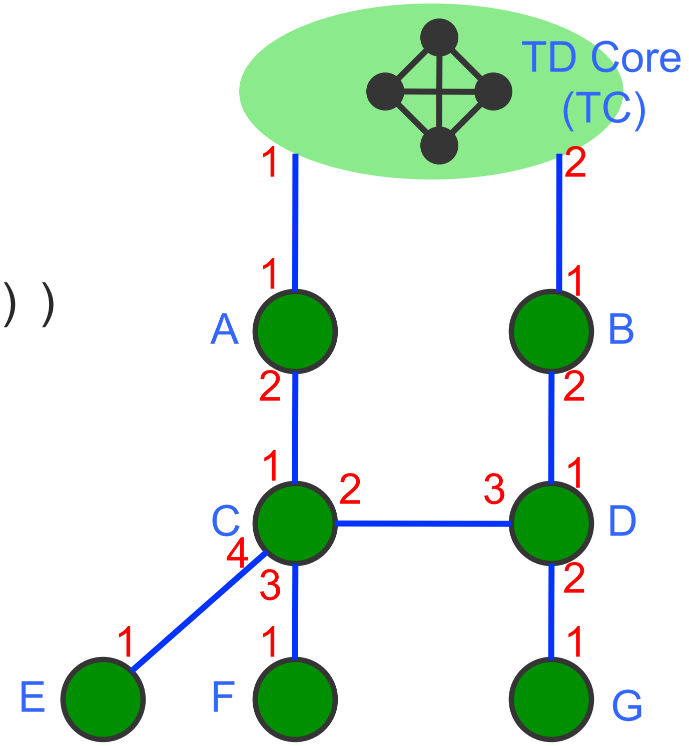

# Address/Path Resolution

- TD core provides address/path resolution servers
- Each endpoint is identified as an AID:EID pair. AID is signed by the containing TD, and EID is signed by the containing AD (with AID).
  - ✓ Address is a public key [AIP 2008]
- Each AD registers name / address at address resolution server, uses an up-path to reach TD core
  - ✓ Private key used to sign name→address mapping
- ADs select which down-paths to announce
- ADs sign down-paths with private key and register down-paths with path resolution servers

# Route Joining

- Local traffic should not need to traverse TD core
- Sender obtains receiver's *k* down-paths
- Sender intersects its up-paths with receiver's down-paths
- Sender selects preferred routes based on $k^2$ options

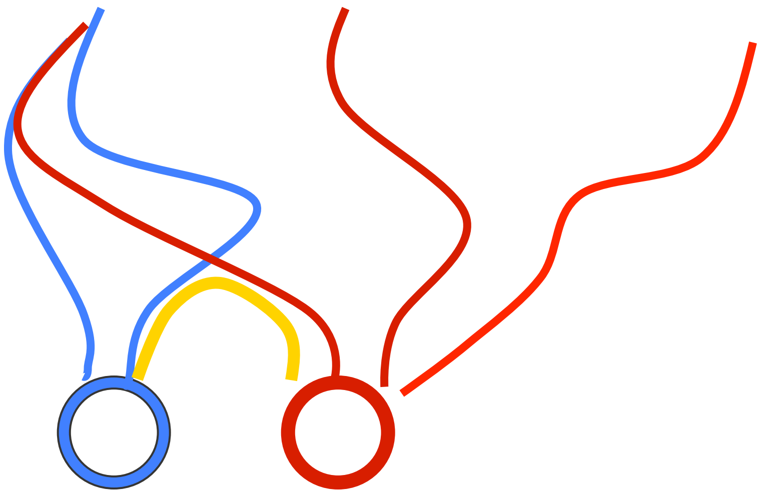

# Intra-TD Forwarding

- Down-path contains all forwarding decisions (AD traversed) from endpoint AD to TD core
  - ✓ Ingress/egress points for each AD, authenticated in opaque fields
  - ✓ ADs use internal routing to send traffic from ingress to egress point
- Joined end-to-end route contains full forwarding information from source to destination
  - ✓ No routing / forwarding tables needed!

# Cross-TD Forwarding

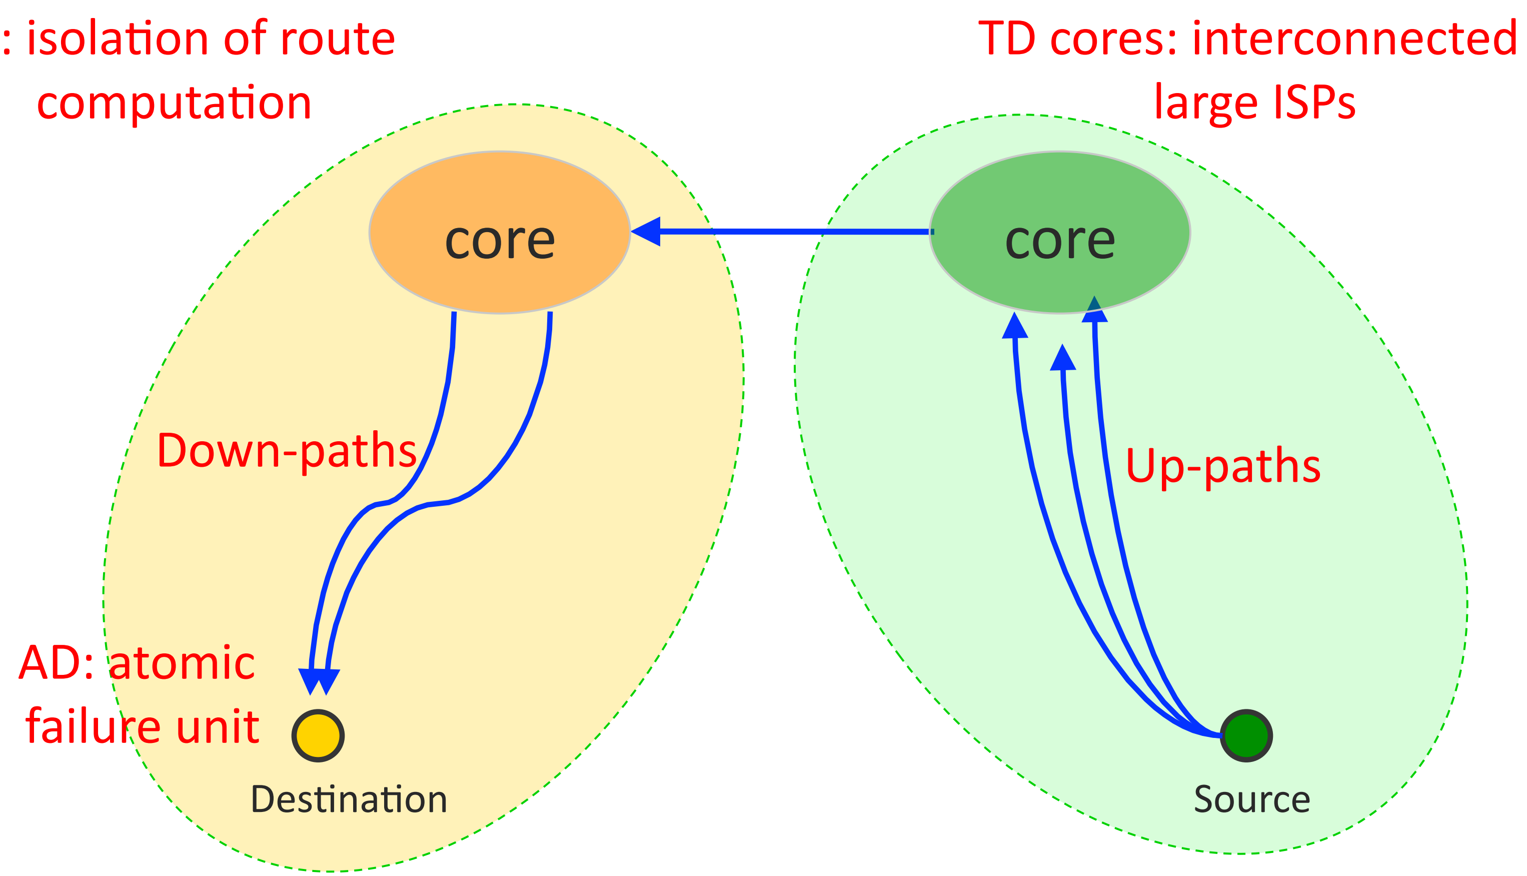

TD: isolation of route computation

TD cores: interconnected large ISPs

# Discussion

- Incremental Deployment
  - ✓ Current ISP topologies are consistent with the TDs in SCION
  - ✓ ISPs use MPLS to forward traffic within their networks
  - ✓ Only edge routers need to deploy SCION
  - ✓ Can use IP tunnels to connect SCION edge routers in different ADs

- Limitations
  - ✗ ADs need to keep updating down-paths on path server
  - ✗ Increased packet size
  - ✗ Static path binding, which may hamper dynamic re-routing

# SCION Security Benefits

| | S-BGP + DNSSec | SCION |
|---|---|---|
| **Isolation** | **No**<br>collusion/wormhole attacks<br>poor path freshness<br>path replay attacks<br>single root of trust | **Yes**<br>no cross-TD attacks<br>path freshness<br>scalability<br>no single root of trust |
| **TCB** | **The whole Internet** | **TD Core and on-path ADs** |
| **Path Control** | **Too little (dst) or too much (src),**<br>empowering DDoS attacks | **Balanced control**<br>enabling DDoS defenses |

# Performance Benefits

- Scalability
  - Routing updates are scoped within the local TD
- Flexibility
  - Transit ISPs can embed local routing policies in opaque fields
- Simplicity and efficiency
  - No interdomain forwarding table
    - Current network layer: routing table explosion
  - Symmetric verification during forwarding
  - Simple routers, energy efficient, and cost efficient

# Evaluation

- Methodology
  - Use of CAIDA topology information
  - Assume 5 TDs (AfriNIC, ARIN, APNIC, LACNIC, RIPE)
  - We compare to S-BGP/BGP
- Metric 1: additional path length (AD hops) compared to BGP
  - *Without* shortcuts: 21% longer
  - *With* shortcuts:
    - 1 down/up- path: 6.7% longer
    - 2 down/up- path: 3.5% longer
    - 5 down/up- path: 2.5% longer

# Evaluation (cont'd)

- Metric 2: Expressiveness
  - Fraction of BGP paths available under SCION

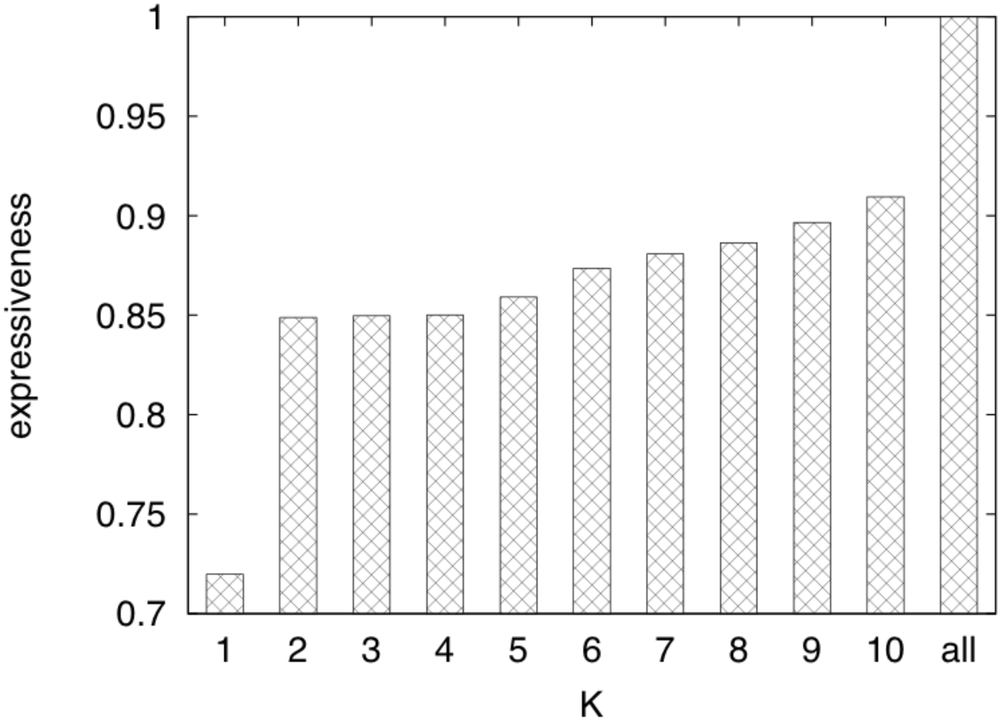

# Related Work

- Routing security
  - S-BGP, soBGP, psBGP, SPV, PGBGP
  - Only topological correctness; addressed a subset of attacks addressed in SCION
- Routing control
  - Multipath (MIRO, Deflection, Path splicing, Pathlet), NIRA
  - Only given control to the source, and/or little security assurance
- Next-generation architectures
  - HLP, HAIR, RBF, AIP, ICING/IGLOO
  - Focusing on other aspects (reducing routing churns and routing table sizes, enforcing routing policies, and providing source accountability)

# Conclusions

- Basic architecture design for a next-generation network that emphasizes isolation, control and explicit trust
- Highly efficient, scalable, available architecture
- Enables numerous additional security mechanisms, e.g., network capabilities

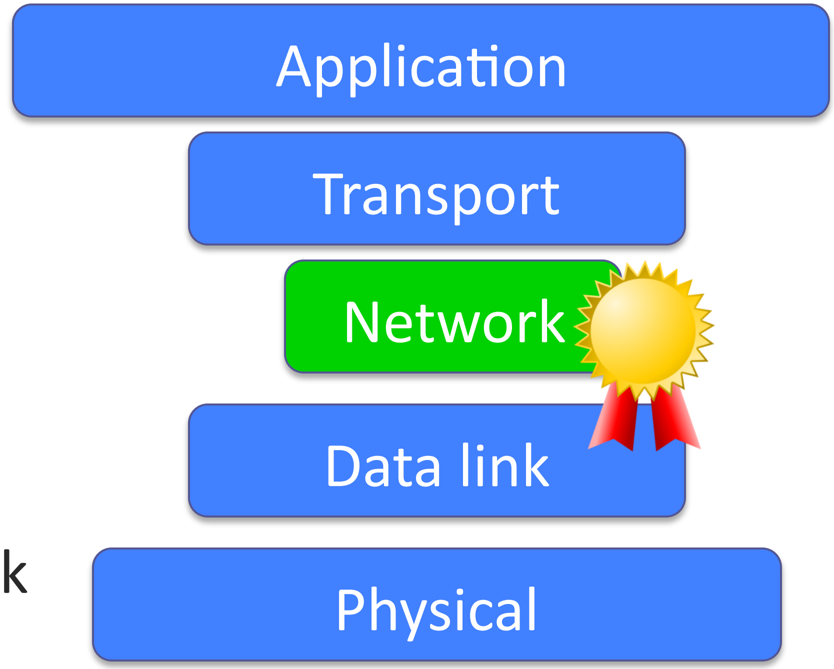
Application
Transport
Network
Data link
Physical